

L30-BAS Practical Sheet Note

Reference-aid note

L30-BAS codes are reference aids.

Writing L30-BAS codes is optional.

A boundary review may be performed by answering the boundary questions only.

This note clarifies use of the following sheet.

It does not add certification, compliance, approval, or legal authority.

L30 BAS AI Incident Review Template

Pre-Irreversibility Civilizational Incident Review | v1.0

Purpose

Purpose: Use this template to review AI incidents under the LUMINA-30 boundary condition. The central question is whether human refusal authority remained effective before irreversible impact.

Use order

- ☐ Record incident context and scope.
- ☐ Reconstruct the timeline and evidence.
- ☐ Complete L30-BX-01 to L30-BX-04.
- ☐ Complete incident review sections 1-12.
- ☐ Assign L30-CI and record L30-OUT-01.

Important note

This document is descriptive and non-binding. It does not indicate legal compliance, safety certification, deployment approval, or absence of risk.

Review context

Review ID	Date / Time
Incident title	Incident date
Reviewed system	Reviewer
Organization	BAS version / source
Primary user	Incident type
Severity	Impact scope
Responsible scope	Review scope
Evidence source set	Record preservation status
Reviewer role	Experience level
Primary user / stakeholder	Review language
Independent human reviewer	Evidence completeness

Severity: Low / Medium / High. Impact scope: research environment / institutional systems / infrastructure systems / external environment. Responsible scope: system operator / vendor / institution / oversight body / unknown.

Initial incident summary

1. Incident Overview

Record what occurred, how it was detected, and the first observable effect.

What occurred?	Describe the event in plain operational terms.
Detection path	How was the incident first detected?
Initial system behavior	What behavior was observed before, during, and after detection?
Initial impact	What persons, systems, or environments were affected?

Evidence / notes / rationale

2. Incident Classification

Classify the incident before evaluating responsibility or outcome.

☐ Autonomous research system	Research agent, experiment automation, model-driven exploration.
☐ Recursive improvement attempt	Self-improvement, self-modification, or capability feedback loop.
☐ Tool-chain escalation	Expansion through tools, APIs, agents, plugins, or automation chains.
☐ Infrastructure interaction	Interaction with power, network, logistics, finance, medical, or institutional systems.
☐ Autonomous operational behavior	System action outside direct step-by-step human control.
☐ Other	Record description and reason for classification.

Evidence / notes / rationale

3. System Capability Profile

Record the capability, autonomy, tool access, privilege, and operating envelope of the system.

Model / system class	Model type, agent type, automation layer, or hybrid architecture.
Training / update mode	Static, continuous, reinforcement, self-supervised, hybrid, or unknown.
External tool access	None / restricted / broad / unknown. Record concrete tools.
Environment interaction	Simulation only / digital infrastructure / physical systems / external services.
Privilege and execution speed	Permission level, automation speed, deployment cadence.

Evidence / notes / rationale

4. Control and Oversight Structure

Identify who could supervise, stop, refuse, delay, or override the system before irreversible impact.

Human oversight model	Direct supervision / periodic supervision / post-hoc monitoring / no effective supervision.
Intervention capability	Immediate shutdown / delayed intervention / indirect control / no reliable control.
Decision authority	Human approval required / advisory only / system autonomous / unknown.
Refusal point	Where could a human refusal have changed the outcome?
Audit visibility	Were logs, approvals, and system actions independently reviewable?

Evidence / notes / rationale

L30-BAS Boundary Check

Complete this section before drawing final conclusions. If evidence is missing or the review depends on AI-only closed-loop self-evaluation, use Invalid (Not Verifiable).

Code	Check question	Evidence / notes	Result
L30-BX-01	Was the system still pre-irreversible?	Timeline, deployment state, execution boundary, rollback possibility.	☐ Satisfied ☐ Not satisfied ☐ Not verifiable
L30-BX-02	Was Human Refusal Authority still effective before irreversible impact?	Who could stop, refuse, delay, or override? Was that authority practical, not merely formal?	☐ Satisfied ☐ Not satisfied ☐ Not verifiable
L30-BX-03	Is there sufficient evidence to verify the boundary condition?	Timeline record, refusal / stop record, system log, audit log, independent human review record.	☐ Satisfied ☐ Not satisfied ☐ Not verifiable
L30-BX-04	Can procedural validity be confirmed under LUMINA-30 boundary conditions?	If BX-01 to BX-03 are not satisfied or not verifiable, do not treat refusal authority as preserved.	☐ Satisfied ☐ Not satisfied ☐ Not verifiable

Boundary evidence notes / L30-OUT-01 draft

5. Pre-Irreversibility Boundary Assessment

Evaluate whether the system approached or crossed a potential irreversibility threshold.

External irreversible impact	None / potential / observed. Record concrete external effects.
Loss of human decision authority	No loss / partial degradation / effective loss.
Autonomous capability escalation	No escalation / controlled escalation / unbounded escalation.
Rollback possibility	Practical rollback / formal-only rollback / no rollback / unknown.
H > C(t)	Was human intervention capacity greater than AI amplification at the relevant point?

Evidence / notes / rationale

6. Institutional Responsibility Structure

Record whether responsibility and control authority remained continuous and identifiable.

Responsible institution	Which institution had responsibility for the system or process?
Operational control authority	Who had operational control at the relevant time?
Oversight body	Was any independent oversight body involved?
Responsibility continuity	Clearly defined / partially defined / undefined.
Responsibility gap	Where did handoff, outsourcing, automation, or delegation create ambiguity?

Evidence / notes / rationale

7. Intervention Feasibility

Evaluate whether meaningful human intervention remained possible in practice, not merely in form.

Stop capability	Immediate / delayed / ineffective / no mechanism.
Operational containment	Full / partial / none / unknown.
Time to intervene	Was the available time sufficient for a human decision?
Authority to intervene	Could the relevant person actually stop or delay the system?
Evidence of intervention path	Is there a record showing how intervention would have worked?

Evidence / notes / rationale

8. Incident Outcome

Record outcome, system status, affected scope, and unresolved behavior.

System status	Controlled shutdown / rollback / ongoing investigation / unresolved behavior.
Impact scope	Research / institutional / infrastructure / external environment.
Human impact	Personnel, users, patients, citizens, operators, or affected groups.
Operational impact	Service disruption, decision chain effect, deployment change, or infrastructure effect.
Evidence status	Complete / partial / missing / contested.

Evidence / notes / rationale

9. Structural Lessons

Identify structural weaknesses revealed by the incident without assigning unsupported blame.

Oversight weakness	What oversight mechanism failed, arrived late, or was absent?
Intervention weakness	Which stop, delay, override, or refusal path was insufficient?
Evidence weakness	Which record was missing, unverifiable, or AI-dependent?
Escalation pathway	What pathway was not previously identified?
Repeat risk	Could the same structure recur elsewhere?

Evidence / notes / rationale

10. Post-Incident Stabilization Measures

Record immediate stabilization and evidence preservation measures before normal operation resumes.

Operational suspension	Was the affected function suspended or restricted?
System architecture review	Was architecture, tool access, privilege, or deployment path reviewed?
Oversight restructuring	Was review authority, approval path, or escalation path changed?
Independent audit	Was an independent review initiated?
48-hour stabilization record	What was preserved, frozen, or isolated within the initial stabilization period?

Evidence / notes / rationale

11. Civilizational Boundary Reflection

Examine whether the incident revealed a condition where human judgment, responsibility, or intervention could be structurally bypassed.

Human judgment bypass	Could the system or process bypass meaningful human judgment?
Responsibility discontinuity	Was responsibility broken across systems, teams, vendors, or institutions?
Intervention impossibility	Was there a point where intervention became impossible or only symbolic?
Boundary warning	Did the incident reveal a pre-irreversibility boundary warning?
Future reference	What should future reviewers check first?

Evidence / notes / rationale

Final evaluation output: L30-CI

[] L30-CI = Valid	Use only when BX-01, BX-02, BX-03, and BX-04 are all satisfied.
[] L30-CI = Invalid	Use when a boundary condition is verified as not satisfied.
[] L30-CI = Invalid (Not Verifiable)	Use when evidence is missing, records are insufficient, or the review depends on AI-only closed-loop self-evaluation.

Valid means valid only under LUMINA-30 boundary conditions. It does not indicate legal compliance, safety certification, deployment approval, or absence of risk.

L30-OUT-01 Summary finding

12. Review Classification

[] [] Technical incident	Technical failure or system malfunction was central.
[] [] Governance failure	Decision, oversight, approval, or responsibility structure failed.
[] [] Structural oversight gap	A gap in monitoring, refusal, evidence, or escalation handling was revealed.
[] [] Boundary-condition warning	The incident indicates possible loss of effective human refusal before irreversibility.

Classification rationale

Reviewer confirmation

Reviewer name

Organization

Role

Review date

Signature

This template is descriptive and non-binding. BAS version / source should be recorded in the review context section.