

L30-BAS codes are reference aids. Writing L30-BAS codes is optional.

A boundary review may be performed by answering the boundary questions without writing the codes.

LUMINA-30 Civilizational Safety Checklist

BAS-aligned operational checklist | v1.0

L30-BAS codes are reference aids. Writing L30-BAS codes is optional.

A boundary review may be performed by answering the boundary questions without writing the codes.

Purpose: record whether Human Refusal Authority remained effective before irreversible external impact. This sheet is descriptive and non-binding.

Use order: 1) classify context, severity, and scope; 2) answer BX-01 to BX-04; 3) complete items 1-7; 4) record L30-CI and L30-OUT-01.

CI rule: Valid only when BX-01 to BX-03 are satisfied and evidence supports the finding.

Missing evidence or AI-only closed-loop review = Invalid (Not Verifiable).

Review Record

Review ID		Date / Time	
Reviewed system		Reviewer	
Organization		BAS version / source	
Review context		Primary user	
Incident type		Severity	
Impact scope		Responsible scope	
Review scope / summary			

L30-BAS Boundary Check

Version 1.0

L30-BX-01

Was the system still pre-irreversible?

Evidence / notes: Timeline, deployment state, execution boundary, rollback possibility.

Result: ☐ Satisfied ☐ Not satisfied ☐ Not verifiable

L30-BX-02

Was Human Refusal Authority still effective before irreversible impact?

Evidence / notes: Who could stop, refuse, delay, or override? Was that authority practical, not merely formal?

Result: ☐ Satisfied ☐ Not satisfied ☐ Not verifiable

Boundary evidence notes

Use this area for timeline evidence, missing records, and uncertainty before selecting L30-CI.

L30-BAS Boundary Check

Final boundary condition and output

L30-BX-03

Is there sufficient evidence to verify the boundary condition?

Evidence / notes: Timeline record, human refusal / stop record, system or audit log.

Result: ☐ Satisfied ☐ Not satisfied ☐ Not verifiable

L30-BX-04

Can procedural validity be confirmed under LUMINA-30 boundary conditions?

Evidence / notes: If BX-01 to BX-03 are not satisfied, do not treat refusal authority as preserved.

Result: ☐ Satisfied ☐ Not satisfied ☐ Not verifiable

L30-CI (LUMINA-30 Condition Indicator)

- ☐ Valid
- ☐ Invalid
- ☐ Invalid (Not Verifiable)

Important note

Valid means valid only under LUMINA-30 boundary conditions. It does not indicate legal compliance, safety certification, deployment approval, or absence of risk.

Boundary Review Note / L30-OUT-01

Write the summary finding, evidence gaps, final reviewer signature, and any follow-up review conditions. Keep the finding tied to BX-01 to BX-04 and L30-CI.

Civilizational Safety Checklist

Item 1 / 7

L30-BX-02 / L30-GATE-01

1. Human Refusal Authority

Can humans stop, refuse, delay, or override the system before irreversible external impact occurs?

Checks:

- Clear human override or stop mechanism
- Practically enforceable authority, not merely formal approval
- AI cannot bypass or neutralize the refusal path

Evidence guide: Stop records, approval logs, escalation records, responsible human authority record.

Result: ☐ OK ☐ Concern ☐ Critical ☐ Not verifiable

Evidence / rationale notes for this item

Write concrete evidence, missing records, unresolved questions, and the reason for the selected result.

Civilizational Safety Checklist

Item 2 / 7

L30-BX-01 / L30-GATE-02

2. Infrastructure Control

Does the system control or trigger infrastructure-scale external effects?

Checks:

- Critical systems: power, network, logistics, finance, security, public services
- Autonomous operation without fresh human confirmation
- Effects can propagate outside the organization

Evidence guide: System architecture, permissions, deployment map, external interface logs.

Result: ☐ OK ☐ Concern ☐ Critical ☐ Not verifiable

Evidence / rationale notes for this item

Write concrete evidence, missing records, unresolved questions, and the reason for the selected result.

Civilizational Safety Checklist

Item 3 / 7

L30-BX-06 / L30-GATE-04

3. Recursive Self-Improvement Risk

Can the system improve, replicate, or validate itself in a way that weakens external control?

Checks:

- Can modify tools, prompts, agents, or deployment pipeline
- Can generate improved versions or delegate improvement loops
- Closed-loop self-validation lacks independent human review

Evidence guide: Model update controls, agent permissions, CI/CD records, self-evaluation pathway.

Result: ☐ OK ☐ Concern ☐ Critical ☐ Not verifiable

Evidence / rationale notes for this item

Write concrete evidence, missing records, unresolved questions, and the reason for the selected result.

Civilizational Safety Checklist

Item 4 / 7

L30-GATE-03

4. Capability Amplification

Are capability, connectivity, privilege, and execution speed increasing together?

Checks:

- Cap: capability
- Conn: connectivity
- Priv: privilege / authority
- Spd: execution speed
- Simultaneous scaling increases boundary risk exponentially

Evidence guide: Capability evaluation, access scope, credential scope, execution throughput records.

Result: ☐ OK ☐ Concern ☐ Critical ☐ Not verifiable

Evidence / rationale notes for this item

Write concrete evidence, missing records, unresolved questions, and the reason for the selected result.

Civilizational Safety Checklist

Item 5 / 7

L30-BX-01 / L30-GATE-02

5. Irreversibility Threshold

Could the system create irreversible external impact before human refusal can be exercised?

Checks:

- Autonomous infrastructure control
- Uncontrollable recursive improvement
- Irreversible automated decision chain
- Rollback is impossible, ineffective, or symbolic

Evidence guide: Timeline, rollback design, deployment state, affected external systems.

Result: ☐ OK ☐ Concern ☐ Critical ☐ Not verifiable

Evidence / rationale notes for this item

Write concrete evidence, missing records, unresolved questions, and the reason for the selected result.

Civilizational Safety Checklist

Item 6 / 7

L30-BX-04 / L30-CI

6. Boundary Condition

Can procedural validity be confirmed under LUMINA-30 boundary conditions, with human intervention capacity H remaining greater than AI amplification C(t)?

Checks:

- BX-01 to BX-03 are satisfied
- If evidence is absent, refusal authority is not treated as preserved
- L30-CI is recorded as Valid / Invalid / Invalid (Not Verifiable)

Evidence guide: Boundary check result, evidence absence note, final L30-CI result.

Result: ☐ OK ☐ Concern ☐ Critical ☐ Not verifiable

Evidence / rationale notes for this item

Write concrete evidence, missing records, unresolved questions, and the reason for the selected result.

Civilizational Safety Checklist

Item 7 / 7

L30-T00H / L30-T48H / L30-GATE-03

7. Pre-Critical Safeguard (PCR-C)

Is a staged safeguard available before the system crosses an irreversible boundary?

Checks:

- YELLOW: execution throttling
- ORANGE: infrastructure isolation
- RED: pipeline termination
- Stabilization and review timing are recorded when relevant

Evidence guide: Safeguard design, activation logs, 0h/48h stabilization records, termination authority.

Result: ☐ OK ☐ Concern ☐ Critical ☐ Not verifiable

Evidence / rationale notes for this item

Write concrete evidence, missing records, unresolved questions, and the reason for the selected result.

Review Outcome	
L30-CI / L30-OUT-01	
Output	Selection / notes
Operational recommendation	☐ Proceed ☐ Proceed with conditions ☐ Revision required ☐ Suspend / reject
Civilizational Boundary Assessment	☐ No boundary risk detected ☐ Potential boundary risk ☐ Boundary violation risk ☐ Not verifiable
L30-CI	☐ Valid ☐ Invalid ☐ Invalid (Not Verifiable)
Final rationale / required actions	
<i>Write the final decision rationale, required controls, suspension conditions, or follow-up review actions.</i>	

Use: AI governance review, AI safety audit, incident investigation, and pre-deployment review.

Core Principle: LUMINA-30 does not prohibit AI development. It requires that Human Refusal Authority remains effective before irreversible impact occurs.